

Problem Statement:

Banks wants to know, whether credit facility will extend to the customer based on the customer data, for this analysis bank is using individual customer geography details, gender, income, industry employment and experience.

Using Chi2_square_test and ANNOVA test and different ML algo, we are trying to suggest the bank

In [34]:

```
#import the required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import chi2_contingency
import stat
from scipy.stats import iqr
```

In [35]:

```
#read the data sets with display the data
df=pd.read_csv(r"D:\clean_dataset.csv")
df
```

Out[35]:

	Gender	Age	Debt	Married	BankCustomer	Industry	Ethnicity	YearsEmployed
0	1	30.83	0.000	1	1	Industrials	White	1.25
1	0	58.67	4.460	1	1	Materials	Black	3.04
2	0	24.50	0.500	1	1	Materials	Black	1.50
3	1	27.83	1.540	1	1	Industrials	White	3.75
4	1	20.17	5.625	1	1	Industrials	White	1.71
...
685	1	21.08	10.085	0	0	Education	Black	1.25
686	0	22.67	0.750	1	1	Energy	White	2.00
687	0	25.25	13.500	0	0	Healthcare	Latino	2.00
688	1	17.92	0.205	1	1	ConsumerStaples	White	0.04
689	1	35.00	3.375	1	1	Energy	Black	8.25

690 rows × 16 columns



In [36]:

```
#check the datasets information  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 690 entries, 0 to 689  
Data columns (total 16 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Gender                690 non-null    int64  
1   Age                   690 non-null    float64  
2   Debt                  690 non-null    float64  
3   Married               690 non-null    int64  
4   BankCustomer          690 non-null    int64  
5   Industry              690 non-null    object  
6   Ethnicity             690 non-null    object  
7   YearsEmployed         690 non-null    float64  
8   PriorDefault          690 non-null    int64  
9   Employed              690 non-null    int64  
10  CreditScore           690 non-null    int64  
11  DriversLicense        690 non-null    int64  
12  Citizen               690 non-null    object  
13  ZipCode               690 non-null    int64  
14  Income                690 non-null    int64  
15  Approved              690 non-null    int64  
dtypes: float64(3), int64(10), object(3)  
memory usage: 86.4+ KB
```

In [37]:

```
#check the any null values in the datasets  
df.isnull().sum()
```

Out[37]:

```
Gender                0  
Age                  0  
Debt                 0  
Married              0  
BankCustomer         0  
Industry             0  
Ethnicity            0  
YearsEmployed        0  
PriorDefault         0  
Employed             0  
CreditScore         0  
DriversLicense       0  
Citizen              0  
ZipCode              0  
Income               0  
Approved             0  
dtype: int64
```

In [38]:

```
#check the mean,mode details
df.describe()
```

Out[38]:

	Gender	Age	Debt	Married	BankCustomer	YearsEmployed	PriorDe
count	690.000000	690.000000	690.000000	690.000000	690.000000	690.000000	690.00
mean	0.695652	31.514116	4.758725	0.760870	0.763768	2.223406	0.52
std	0.460464	11.860245	4.978163	0.426862	0.425074	3.346513	0.49
min	0.000000	13.750000	0.000000	0.000000	0.000000	0.000000	0.00
25%	0.000000	22.670000	1.000000	1.000000	1.000000	0.165000	0.00
50%	1.000000	28.460000	2.750000	1.000000	1.000000	1.000000	1.00
75%	1.000000	37.707500	7.207500	1.000000	1.000000	2.625000	1.00
max	1.000000	80.250000	28.000000	1.000000	1.000000	28.500000	1.00

In [39]:

```
#check the value counts for each parameters
for i in df:
    print(df[i].value_counts())
```

```
CommunicationServices    38
Utilities                 38
Real Estate              30
Education                25
Research                 10
Transport                 3
Name: Industry, dtype: int64
White                    408
Black                   138
Asian                    59
Latino                   57
Other                    28
Name: Ethnicity, dtype: int64
0.000    70
0.250    35
0.040    33
1.000    31
0.125    30
..
4.165     1
```

First will select the categorical type data for chi2_square_test, here categorical type means "object type", "yes/no" lets check the categorical data and create the separate list.

In [40]:

```
#create the list of categorical type variables
category=['Gender', 'Married', 'BankCustomer', 'Industry', 'Ethnicity', 'PriorDefault', 'Employed', 'DriversLicense', 'Citizen', 'ZipCode', 'Approved']
```

In [41]:

```
df.loc[:,category]
```

Out[41]:

	Gender	Married	BankCustomer	Industry	Ethnicity	PriorDefault	Employed	Drive
0	1	1	1	Industrials	White	1	1	
1	0	1	1	Materials	Black	1	1	
2	0	1	1	Materials	Black	1	0	
3	1	1	1	Industrials	White	1	1	
4	1	1	1	Industrials	White	1	0	
...
685	1	0	0	Education	Black	0	0	
686	0	1	1	Energy	White	0	1	
687	0	0	0	Healthcare	Latino	0	1	
688	1	1	1	ConsumerStaples	White	0	0	
689	1	1	1	Energy	Black	0	0	

690 rows × 11 columns



Statistical Testing using Chisquare:

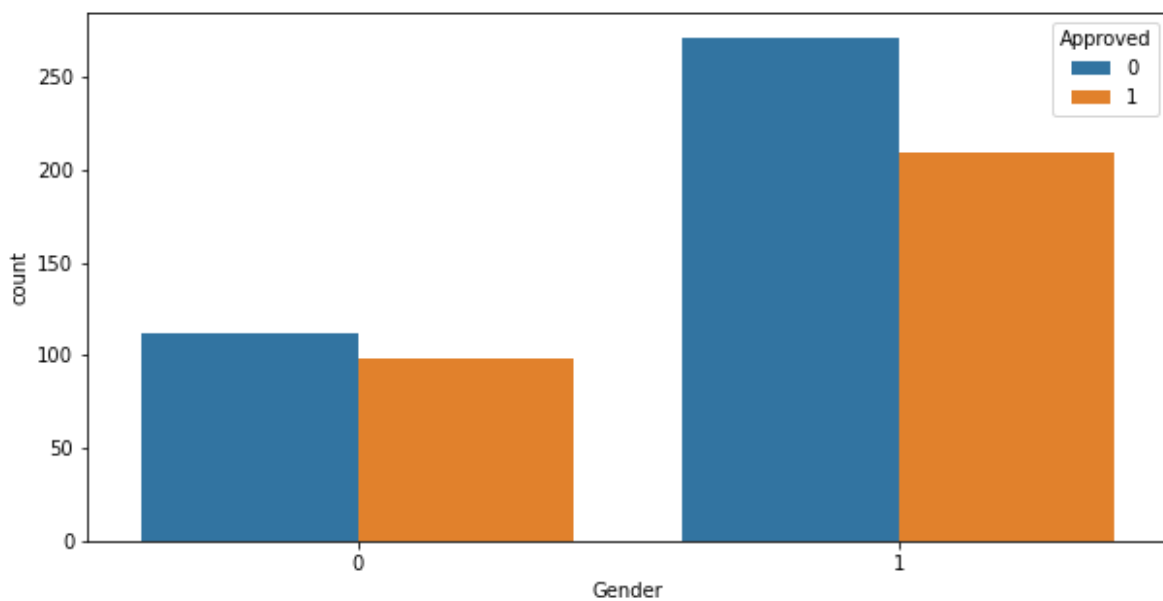
In [42]:

```
import colorama
from colorama import Fore
```

In [43]:

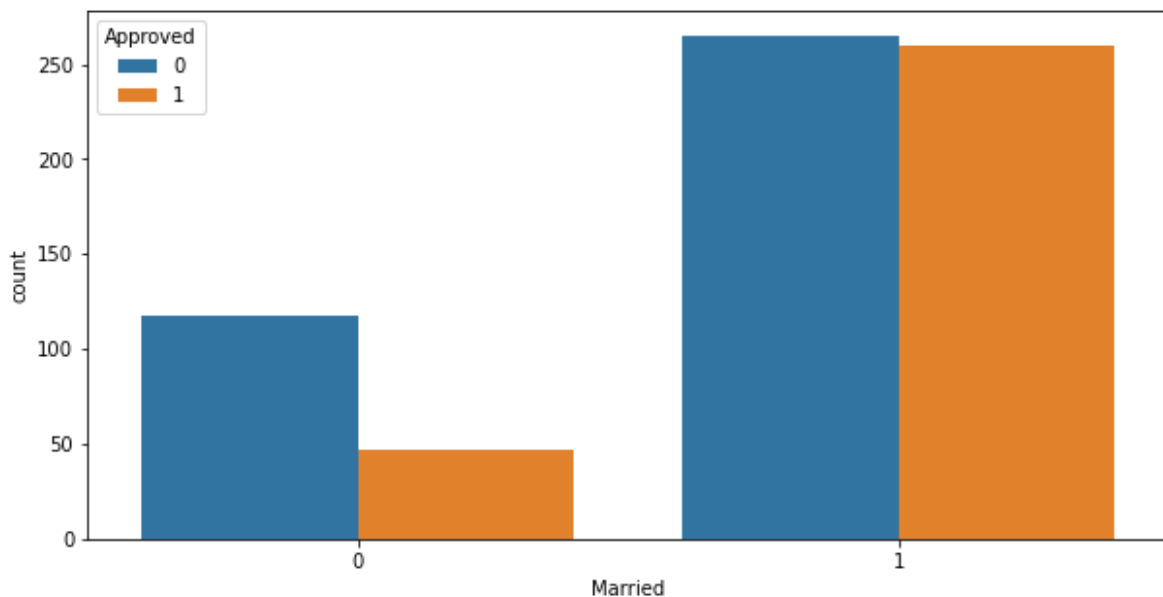
```
for i in category:
    print(i+":")
    plt.figure(figsize=(10,5))
    sns.countplot(x=i,data=df,hue="Approved")
    plt.show()
    a=np.array(pd.crosstab(df.Approved,df[i]))
    (stats,p,dof,_)=chi2_contingency(a,correction=False)
    if p>0.05:
        print(Fore.RED + "'{}'is a 'bad Predictor'".format(i))
        print('p_val={}\\n'.format(p))
    else:
        print(Fore.GREEN + "'{}' is a Good Predictor".format(i))
        print('p_val={}\\n'.format(p))
```

Gender:



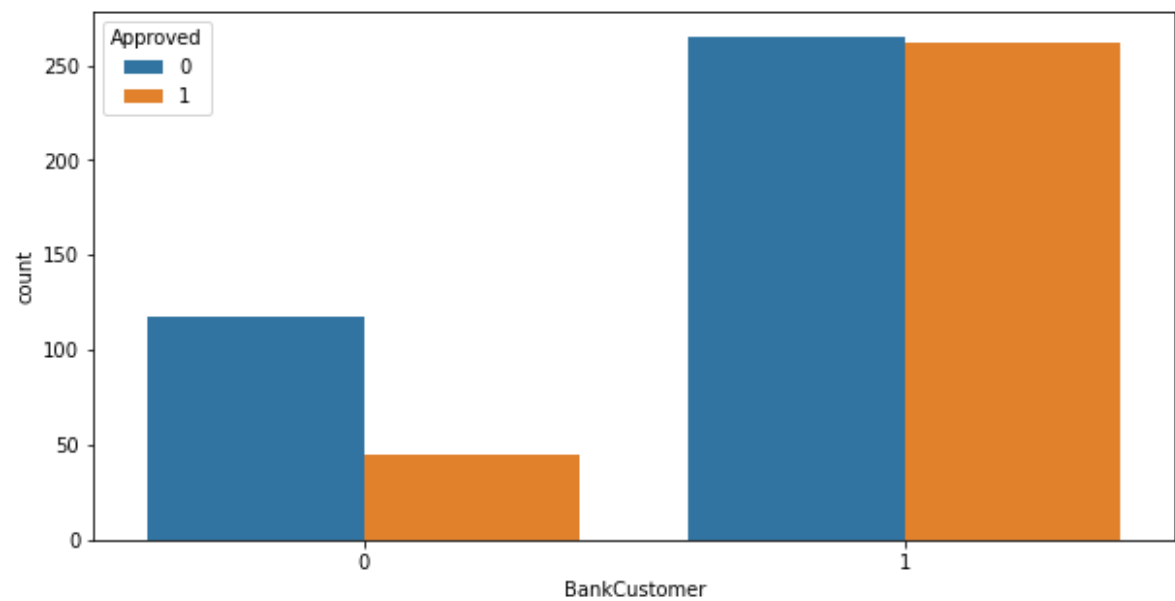
'Gender'is a 'bad Predictor'
p_val=0.44723087514133186

Married:



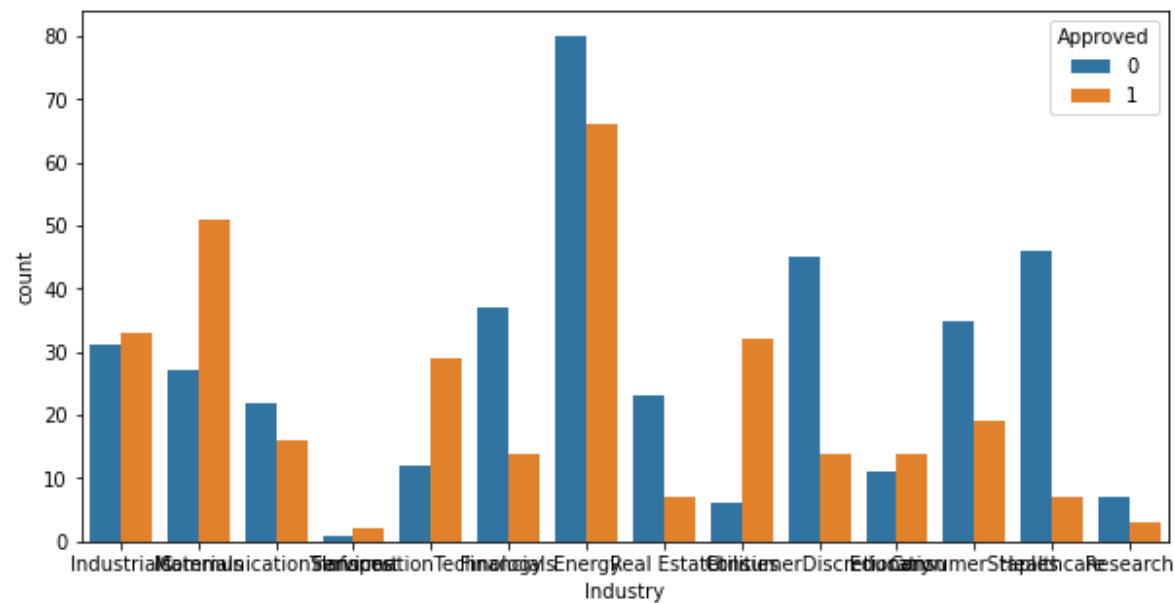
'Married' is a Good Predictor
p_val=2.100231920165588e-06

BankCustomer:



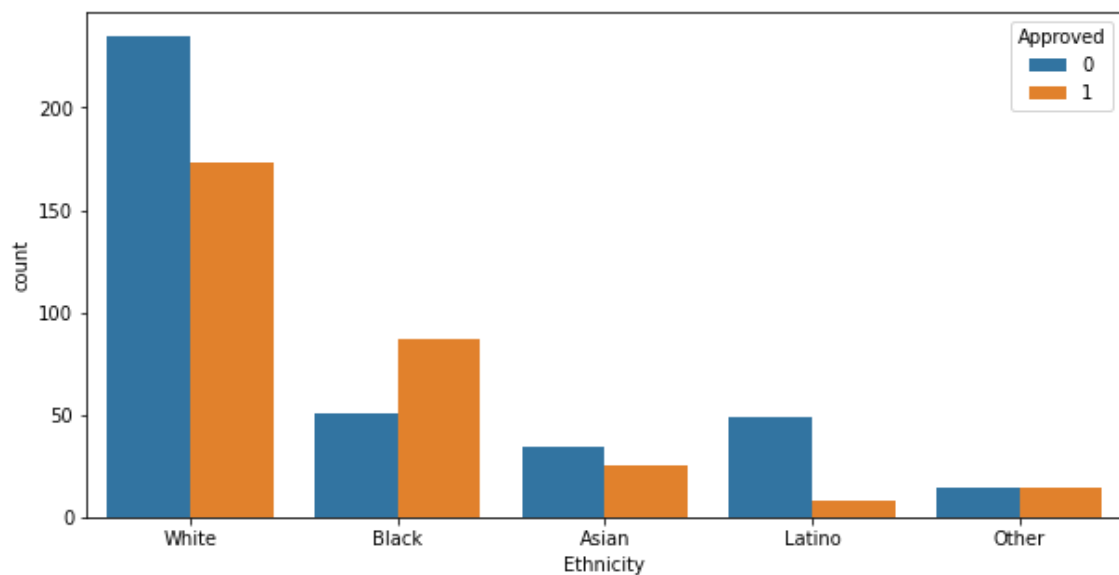
'BankCustomer' is a Good Predictor
p_val=6.91661320541803e-07

Industry:



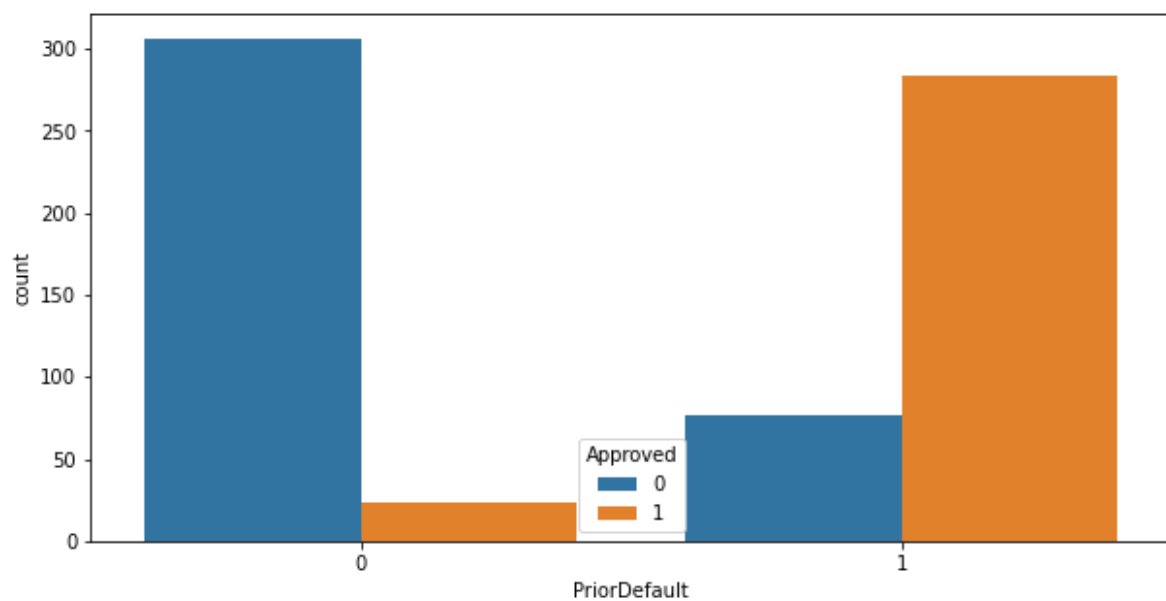
'Industry' is a Good Predictor
p_val=3.502987066102042e-15

Ethnicity:



'Ethnicity' is a Good Predictor
 $p_val=1.823665654934685e-08$

PriorDefault:




'PriorDefault' is a Good Predictor
 $p_val=7.298530125411298e-80$

Employed:




'Employed' is a Good Predictor
 $p_val=2.227269345312281e-33$

DriversLicense:



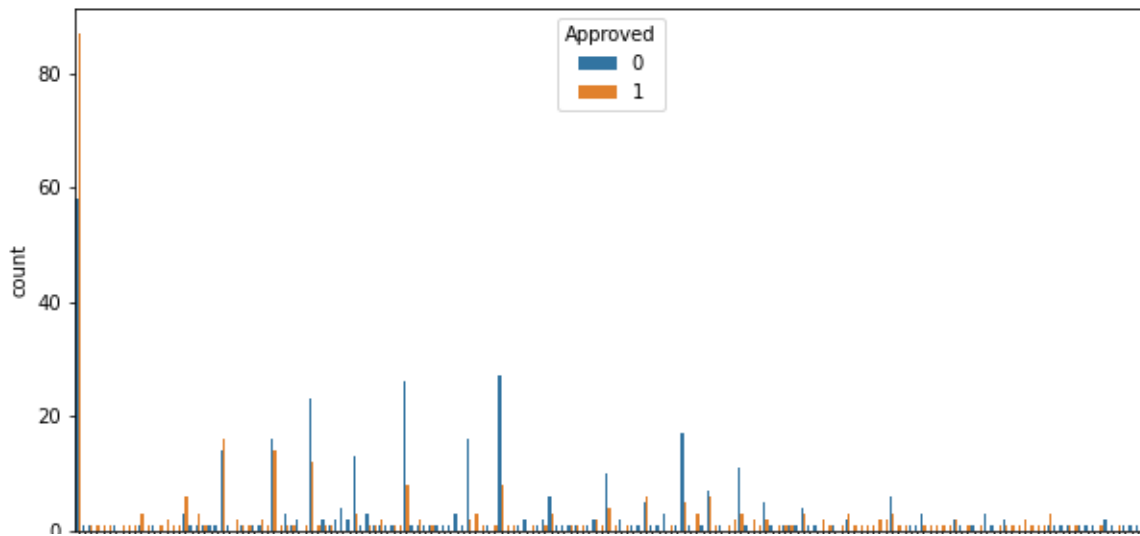
'DriversLicense' is a 'bad Predictor'
 $p_val=0.4061341323141693$

Citizen:



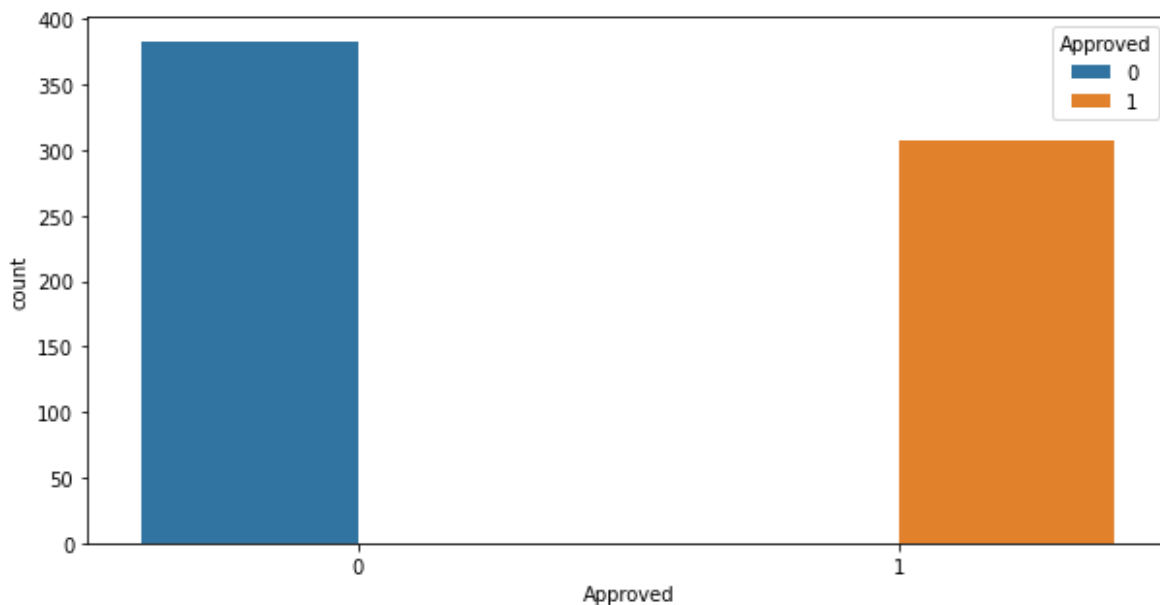
'Citizen' is a Good Predictor
 $p_val=0.010094291370456362$

ZipCode:



'ZipCode' is a Good Predictor
p_val=0.006354824252183887

Approved:



'Approved' is a Good Predictor
p_val=4.469841378183071e-152

Gender:

"we see that majority of credit card holders are males " "since Gender is a bad predictor towards Approval we drop the attribute"

In [44]:

```
df.drop(['Gender'],1,inplace=True)
```

C:\Users\SMFL-20531\AppData\Local\Temp\ipykernel_14400\2126209942.py:1: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.
df.drop(['Gender'],1,inplace=True)

Married:

"we see that majority of credit card holders are Married "

BankCustomer:

"we see that majority of credit card holders are BankCustomers only " "we see that BankCustomers have equal probability of getting approval to not getting approval of credit card " "If the Applicant is not BankCustomers then he/she has higher probability of getting approval of credit card "

Industry:

"we see that majority of credit card holders are people working in Energy sector " "There is High Probability of credit card approval if he/she working in HealthCare sector"

Ethnicity:

"we see that people belonging to White ethnicity use credit card majorly "

PriorDefault:

"we see that people with no defaults have high probability of creditcard approval " "we see that people with payment defaults have less probability of creditcard approval "

DriversLicense:

"we see that majority of credit card holders have DrivingLicense" "since DriversLicense is a bad predictor towards Approval we drop the attribute"

In [45]:

```
df.drop(['DriversLicense'],1,inplace=True)
```

C:\Users\SMFL-20531\AppData\Local\Temp\ipykernel_14400\2114335881.py:1: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```
df.drop(['DriversLicense'],1,inplace=True)
```

Citizen:

"we see that majority of credit card holders have Citizenship bybirth"

Continuous type of variable:

In [46]:

```
column_names=df.columns.tolist()
continious=list(set(column_names)-set(category))
```

In [47]:

```
continious
```

Out[47]:

```
['Income', 'Age', 'YearsEmployed', 'CreditScore', 'Debt']
```

In [48]:

```
df.loc[:,continious]
```

Out[48]:

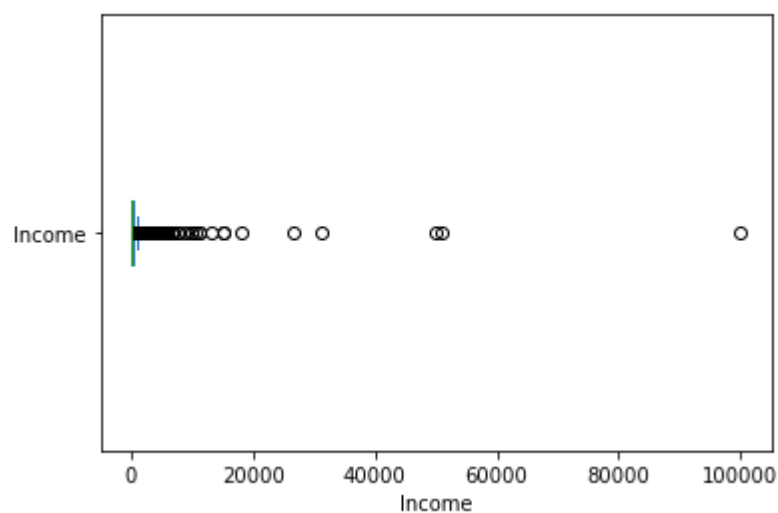
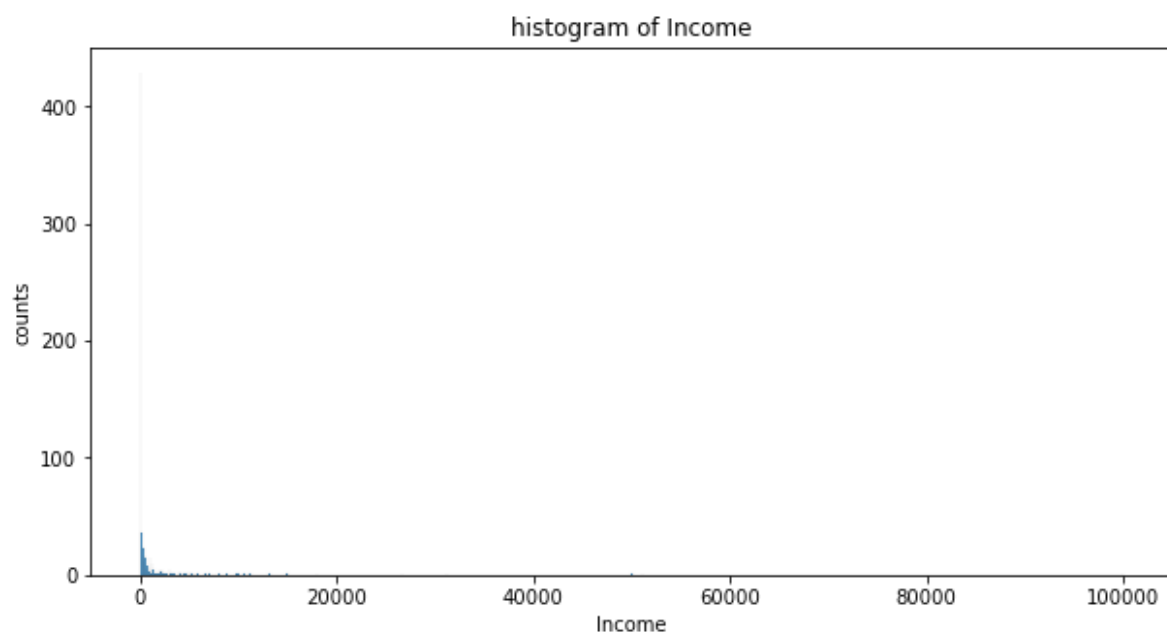
	Income	Age	YearsEmployed	CreditScore	Debt
0	0	30.83	1.25	1	0.000
1	560	58.67	3.04	6	4.460
2	824	24.50	1.50	0	0.500
3	3	27.83	3.75	5	1.540
4	0	20.17	1.71	0	5.625
...
685	0	21.08	1.25	0	10.085
686	394	22.67	2.00	2	0.750
687	1	25.25	2.00	1	13.500
688	750	17.92	0.04	0	0.205
689	0	35.00	8.29	0	3.375

690 rows × 5 columns

In [49]:

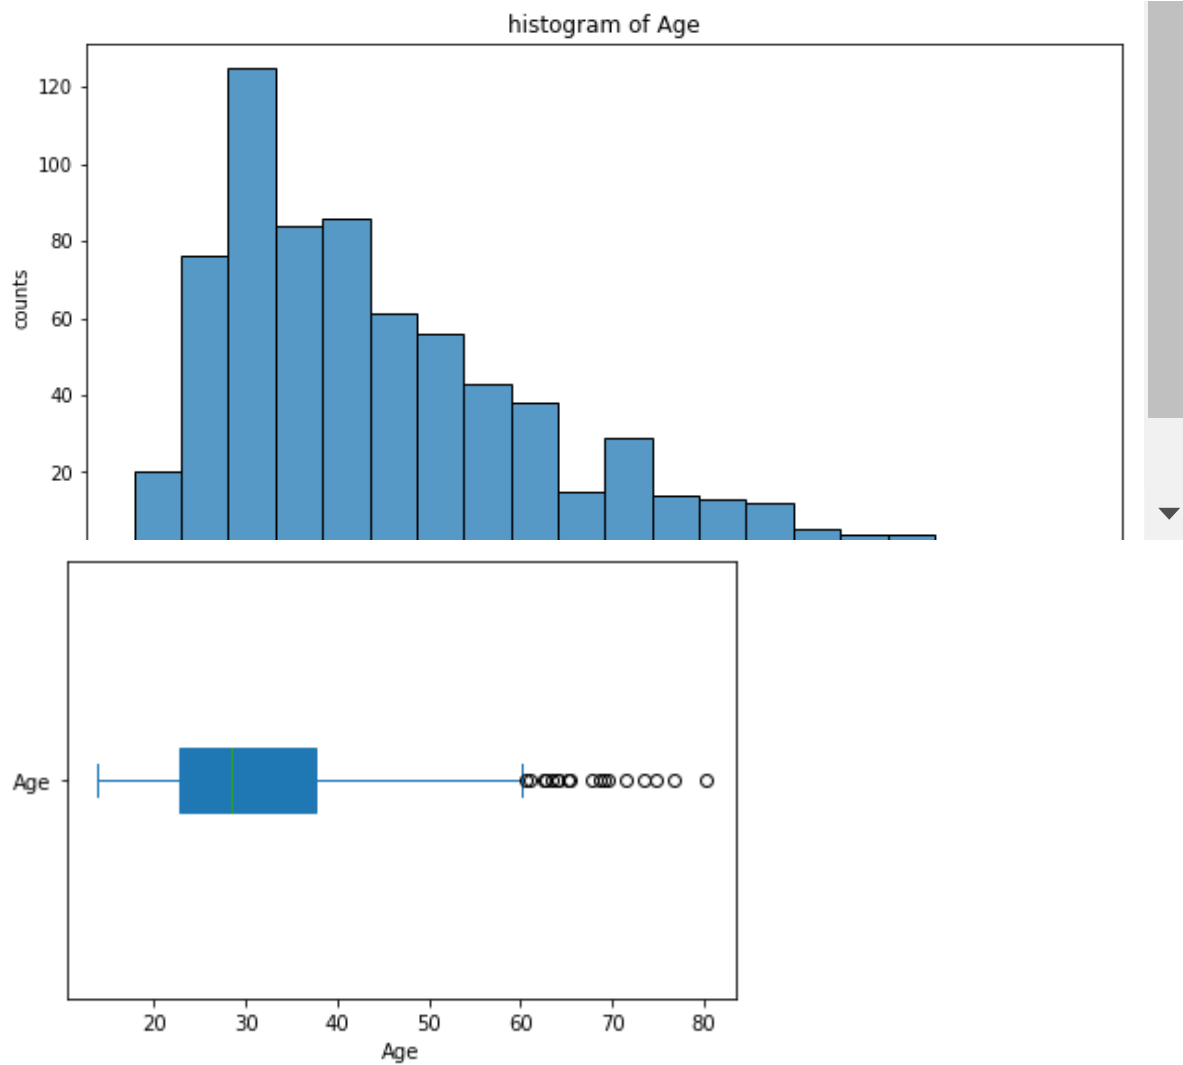
```
for i in continious:
    print(i+':')
    plt.figure(figsize=(10,5))
    sns.histplot(df[i])
    plt.xlabel(i)
    plt.ylabel('counts')
    plt.title('histogram of ' + i)
    plt.show()
    df[i].plot.box(vert=False,patch_artist=True)
    plt.xlabel(i)
    plt.show()
```

Income:

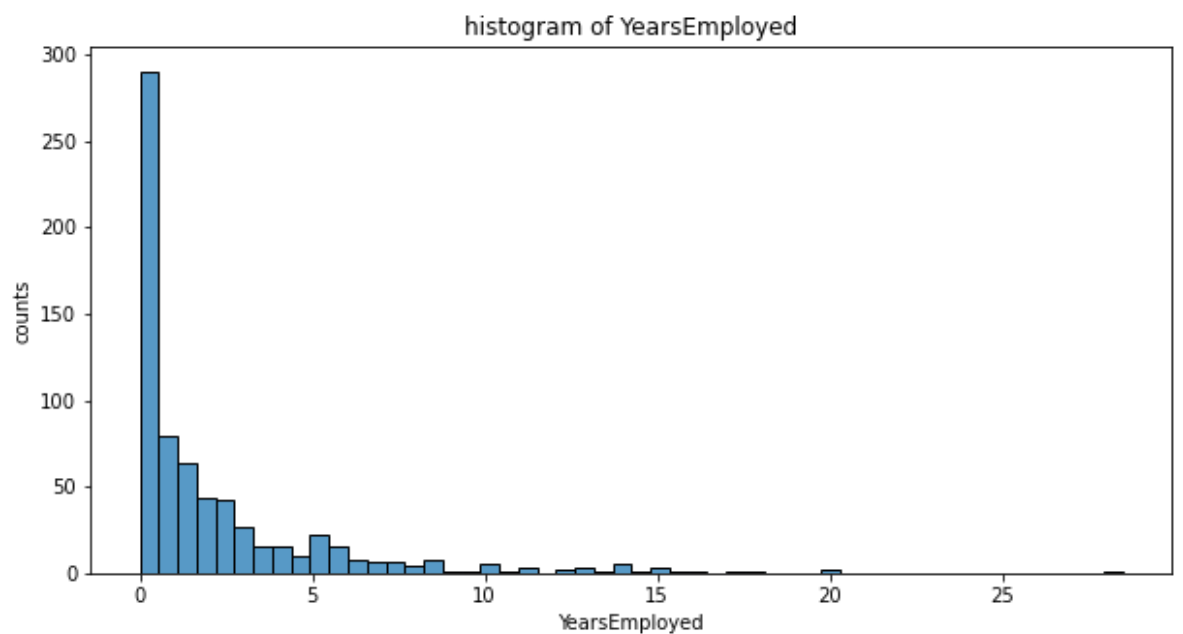



Age:







YearsEmployed:







CreditScore:



Debt:



Treating Outliers:

In [50]:

```
iqr=iqr
```


In [51]:

```
for i in continious:
    print(i+":")
    plt.figure(figsize=(20,10))
    df[i].plot.box(vert=False,patch_artist=True)
    plt.xlabel(i)
    plt.show()

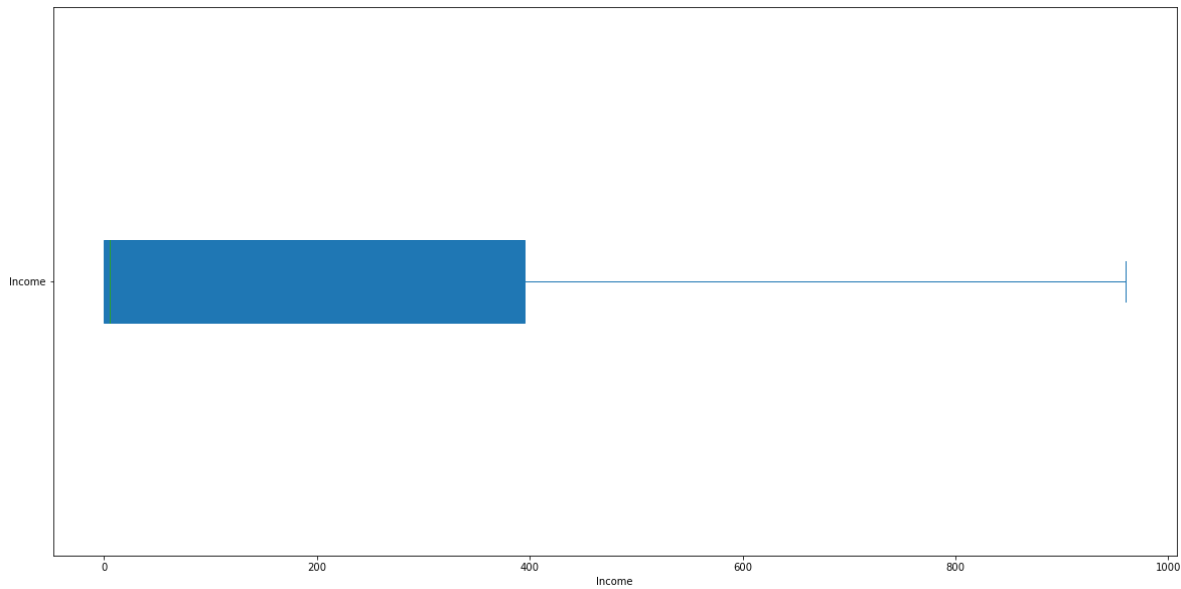
    iqr=df[i].describe()['75%']-df[i].describe()['25%']
    right_lim=df[i].quantile(0.75)+1.5*iqr
    left_lim=df[i].quantile(0.25)-1.5*iqr
    print('Left Limit:{}'.format(left_lim))
    print('Right Limit:{}'.format(right_lim))
    print('iqr:{}'.format(iqr))

a=input("if left outlier type 'left' else 'right' else 'both' if no outliers then type 'no'")
a=a.lower()
if a=='left':
    lo=float(input('by what number do you want to replace'))
    df[i]=[lo if val<left_lim else val for val in df[i]]
elif a=='right':
    ro=float(input('by what number do you want to replace'))
    df[i]=[ro if val>right_lim else val for val in df[i]]
elif a=='both':
    lo=float(input('by what number do you want to replace'))
    ro=float(input('by what number do you want to replace'))
    df[i]=[ro if val>right_lim else val for val in df[i]]
    df[i]=[lo if val<left_lim else val for val in df[i]]
elif a=='no':
    pass
plt.figure(figsize=(20,10))
df[i].plot.box(vert=False,patch_artist=True)
plt.xlabel(i)
plt.show()
```

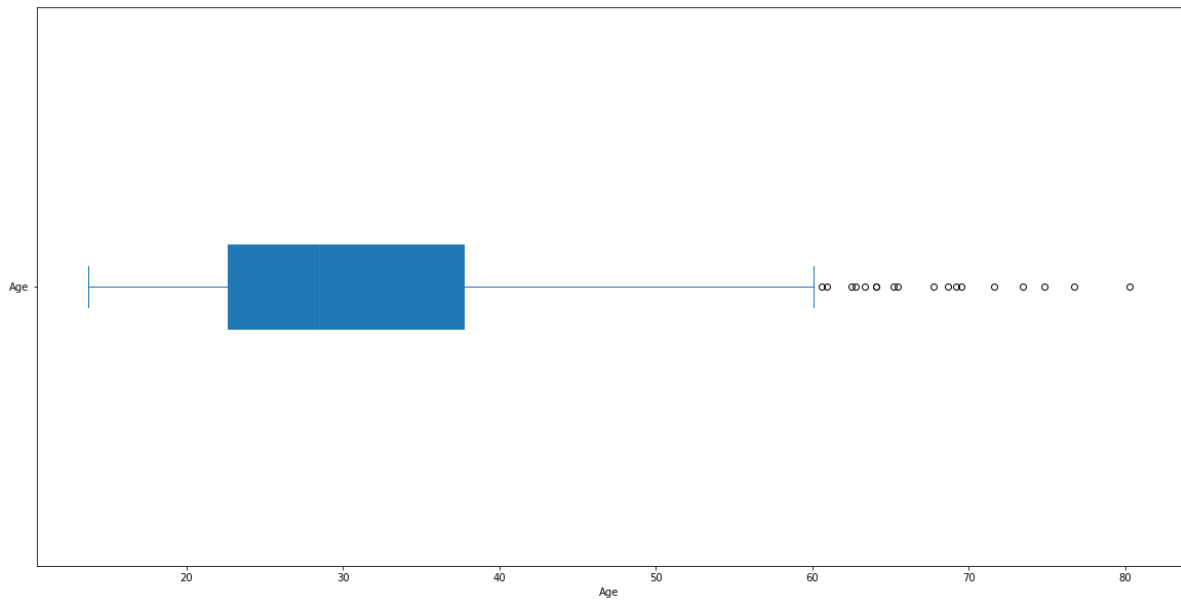
Income:



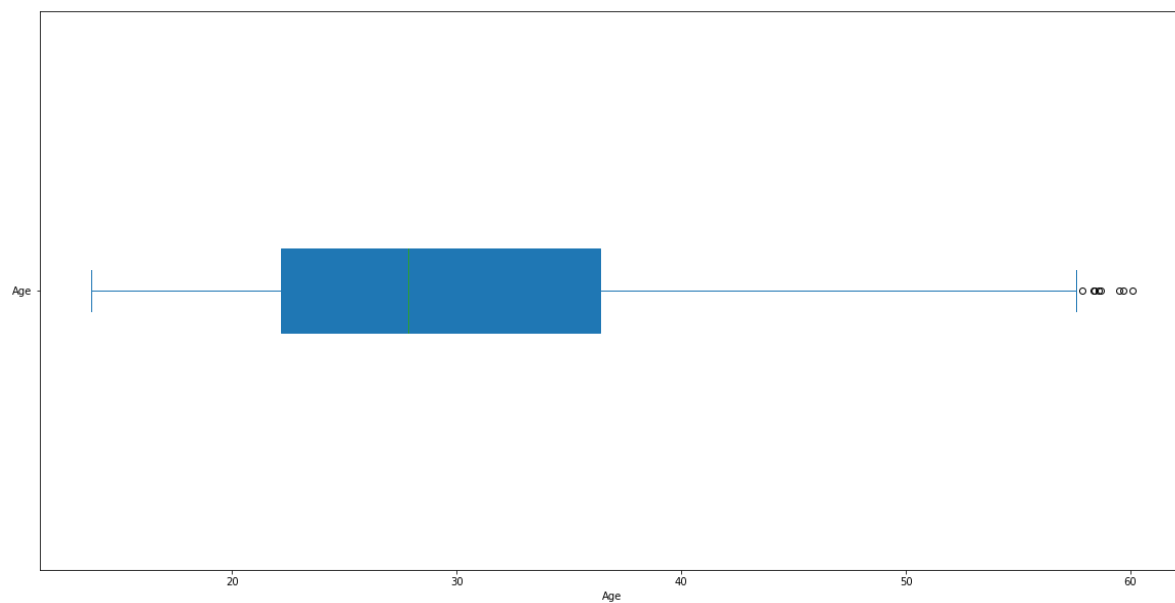
Left Limit:-593.25
Right Limit:988.75
iqr:395.5
if left outlier type 'left' else 'right' else 'both' if no outliers then type 'no'right
by what number do you want to replace395.5



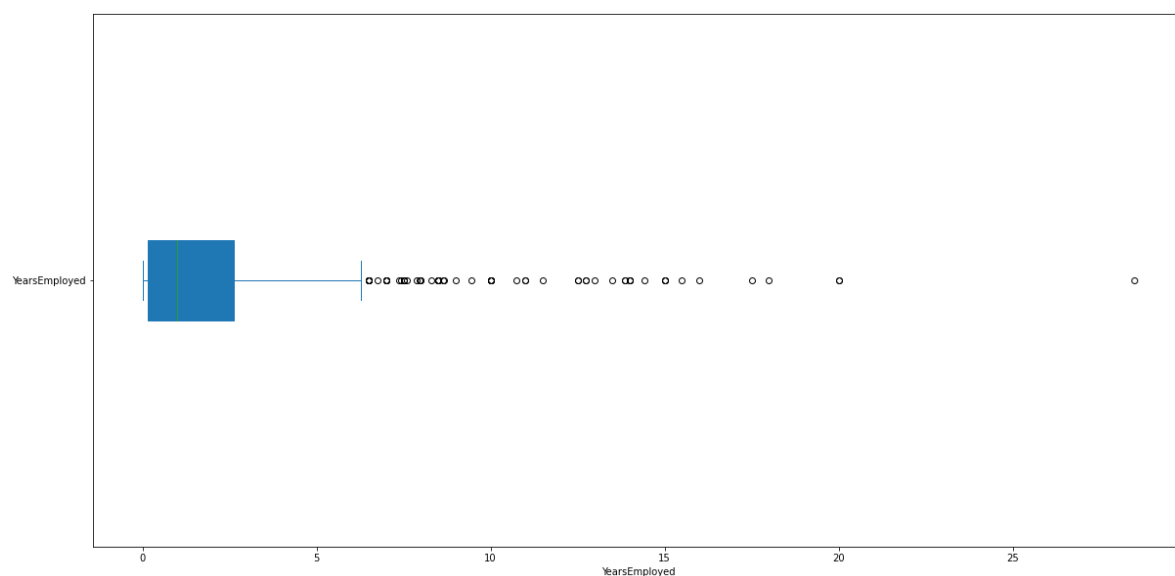
Age:



Left Limit:0.11375000000001023
Right Limit:60.263749999999999
iqr:15.037499999999994
if left outlier type 'left' else 'right' else 'both' if no outliers then type 'no'right
by what number do you want to replace15.03



YearsEmployed:



Left Limit:-3.525

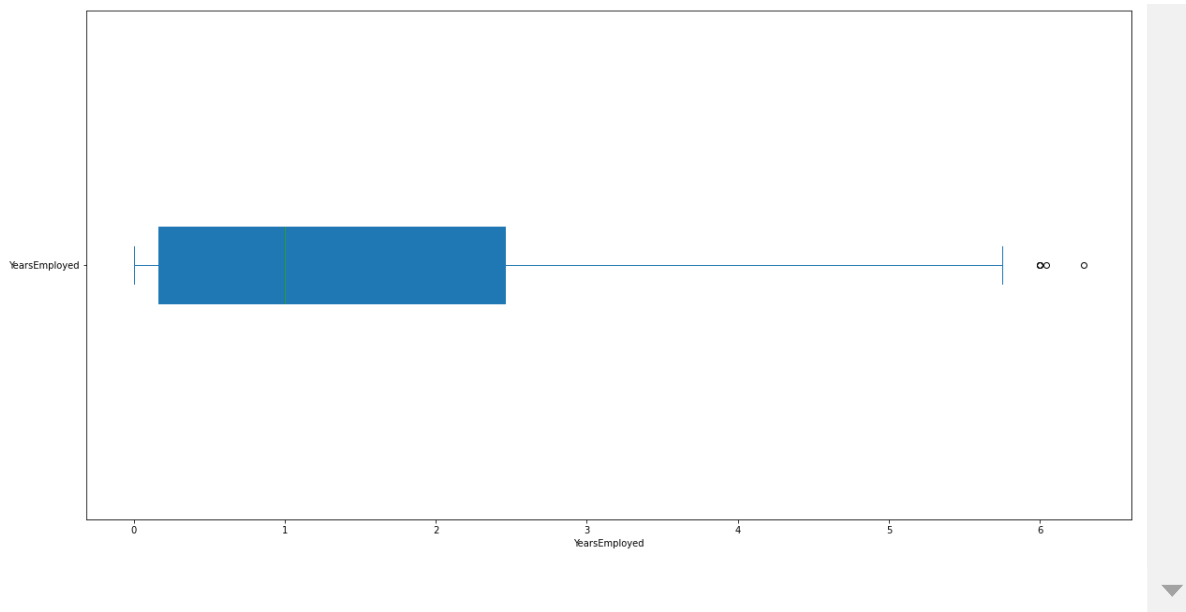
Right Limit:6.314999999999995

iqr:2.46

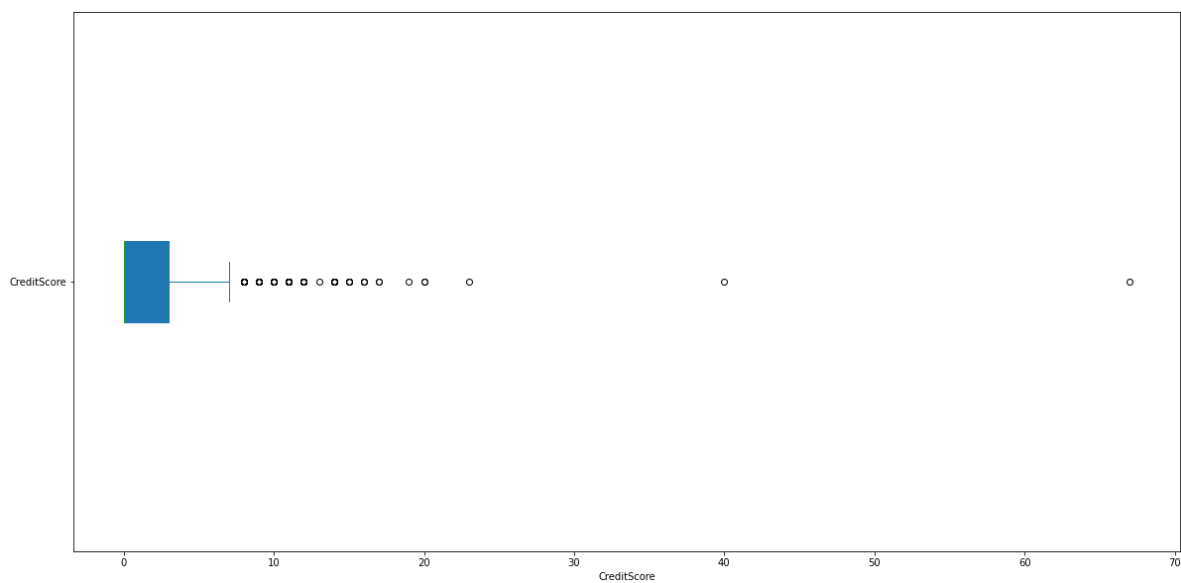
if left outlier type 'left' else 'right' else 'both' if no outliers then type 'no'

by what number do you want to replace 2.46





CreditScore:



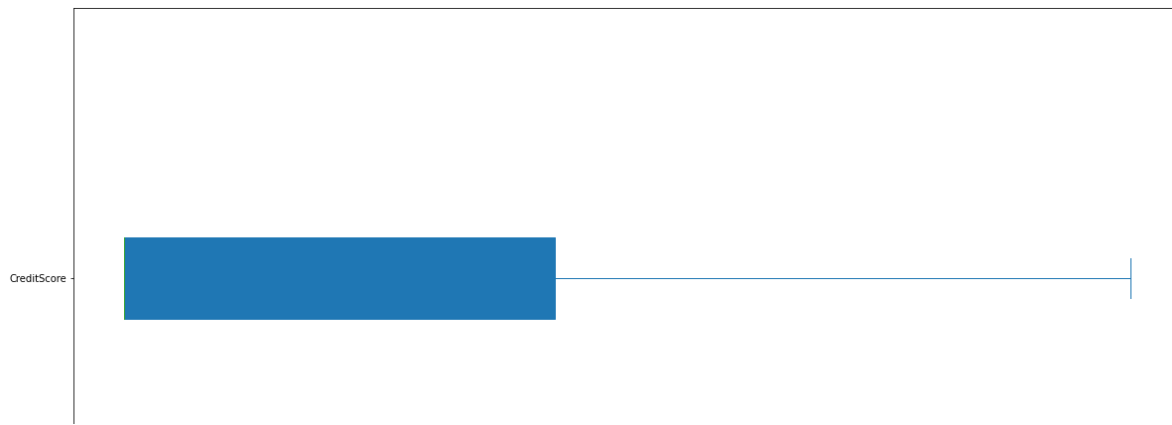
Left Limit:-4.5

Right Limit:7.5

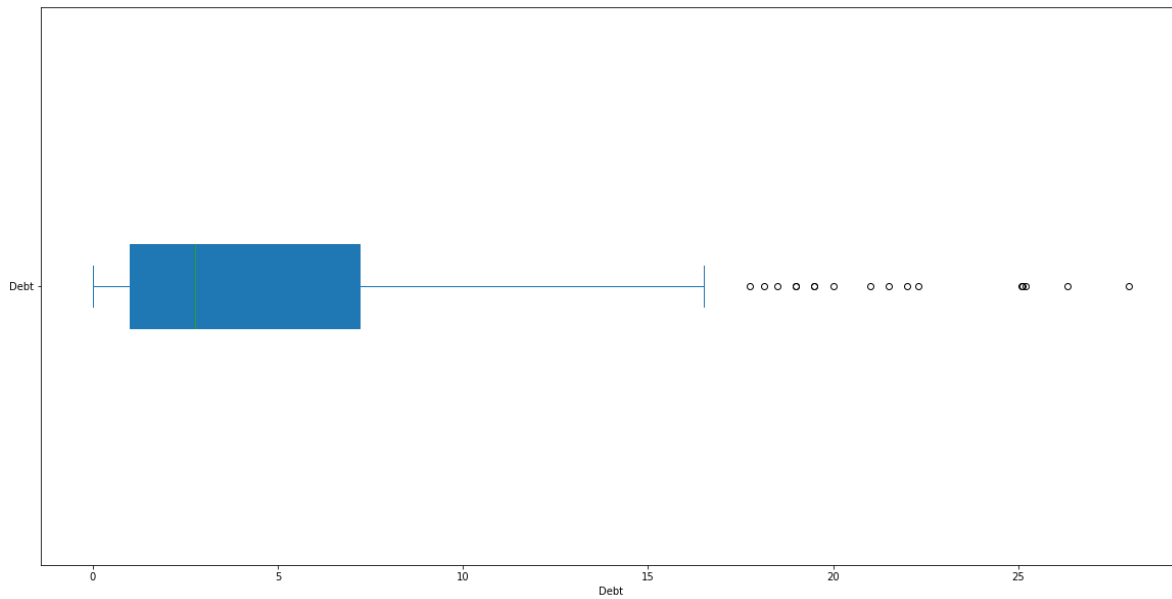
iqr:3.0

if left outlier type 'left' else 'right' else 'both' if no outliers then type 'no'

by what number do you want to replace3.0



Debt:



Left Limit:-8.31125

Right Limit:16.518749999999997

iqr:6.2075

if left outlier type 'left' else 'right' else 'both' if no outliers then type 'no'



by what number do you want to replace 6.2075



In [53]:


```
for i in continious:
    print(i+":")
    plt.figure(figsize=(10,5))
    sns.histplot(df[i])
    plt.xlabel(i)
    plt.ylabel('counts')
    plt.title('histogram of '+i)
    plt.show()
    df[i].plot.box(vert=False,patch_artist=True)
    plt.xlabel(i)
    plt.show()
```

Income:

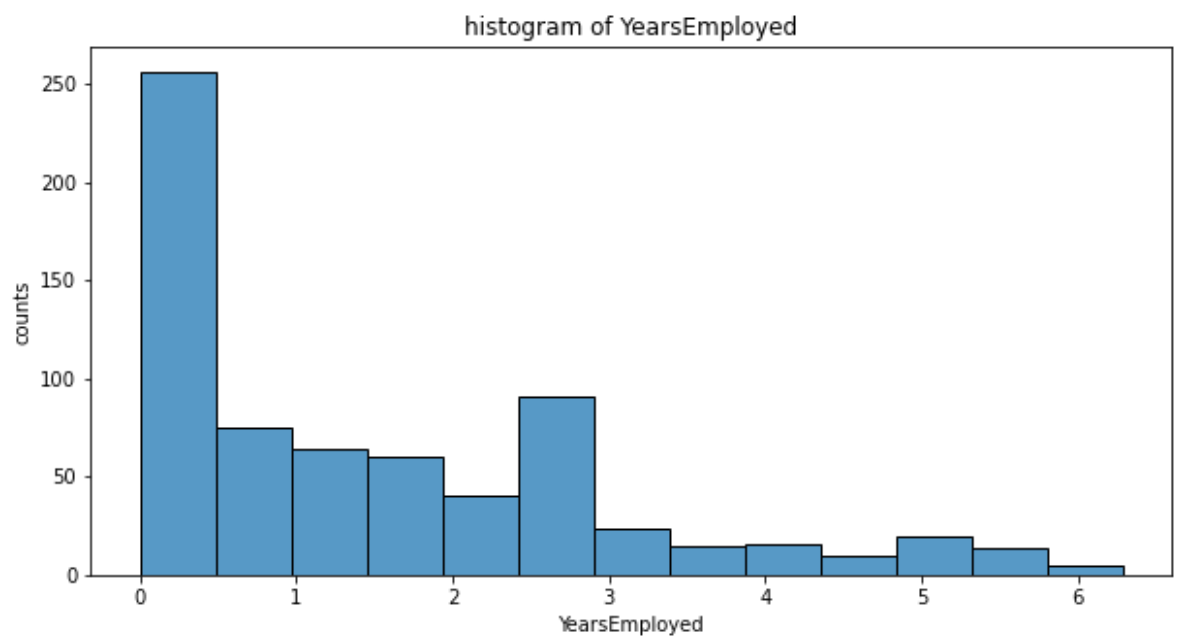


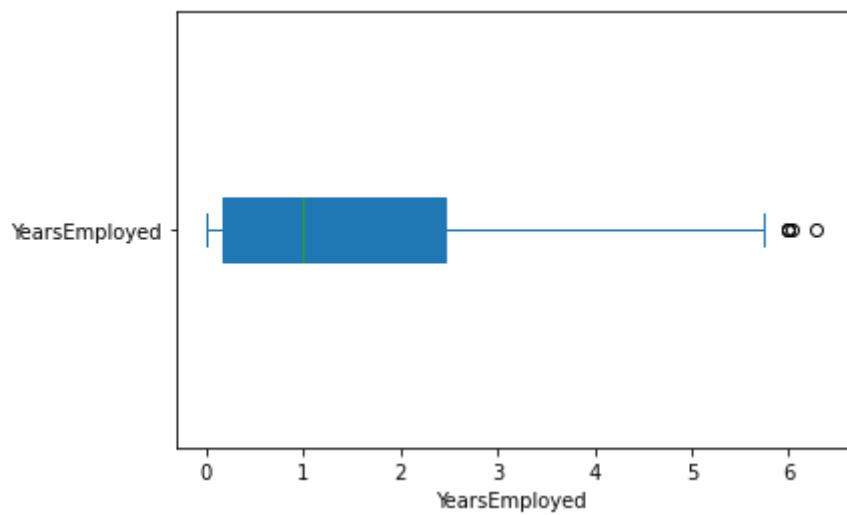
Age:




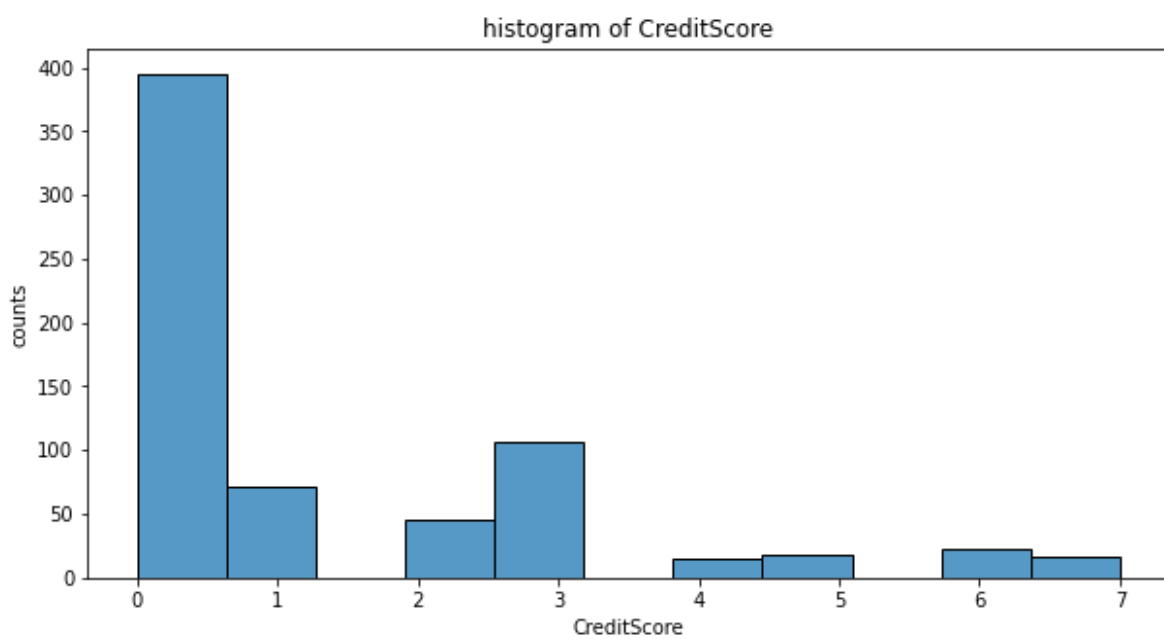


YearsEmployed:






CreditScore:



Debt:

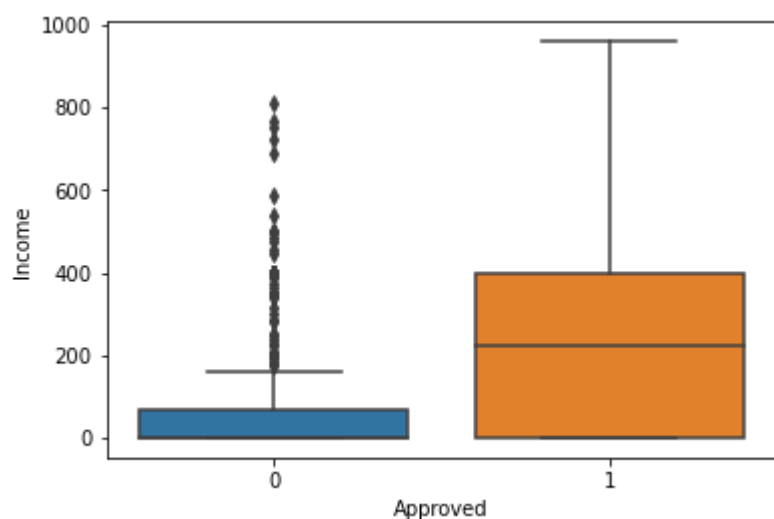
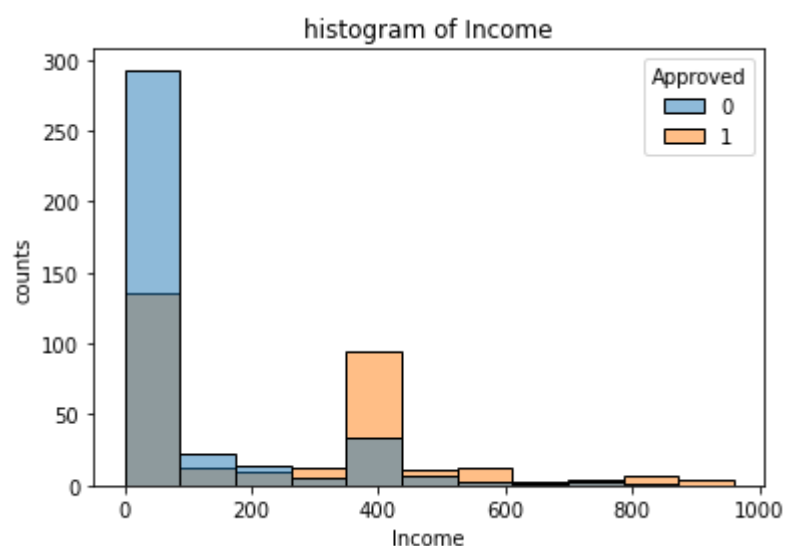
histogram of Debt



In [54]:

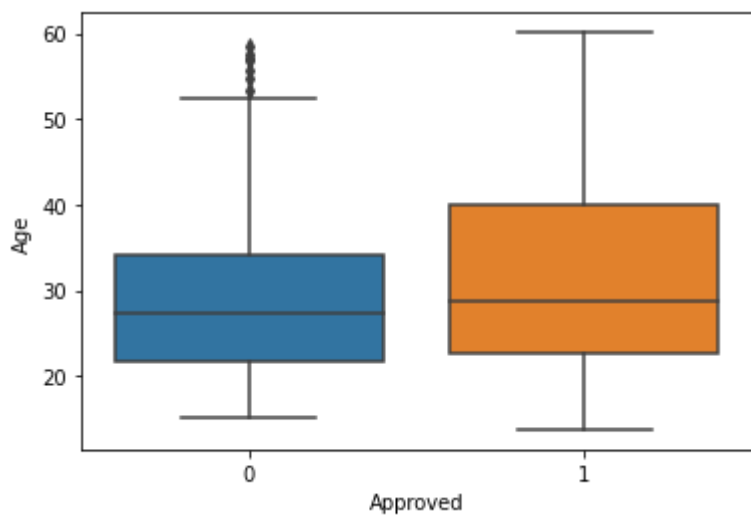
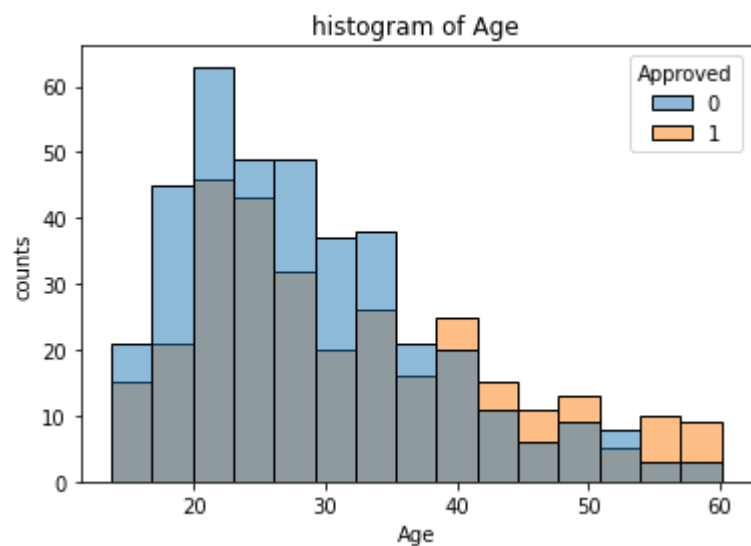
```
for i in continious:
    print(i+":")
    sns.histplot(x=df[i],hue=df.Approved)
    plt.xlabel(i)
    plt.ylabel('counts')
    plt.title('histogram of '+ i)
    plt.show()
    sns.boxplot(y=df[i],x=df.Approved)
    plt.show()
```

Income:

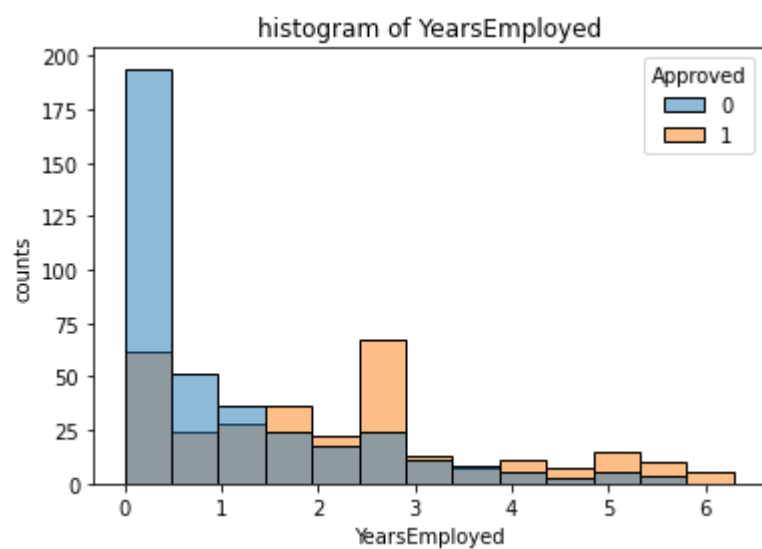


Age:







YearsEmployed:

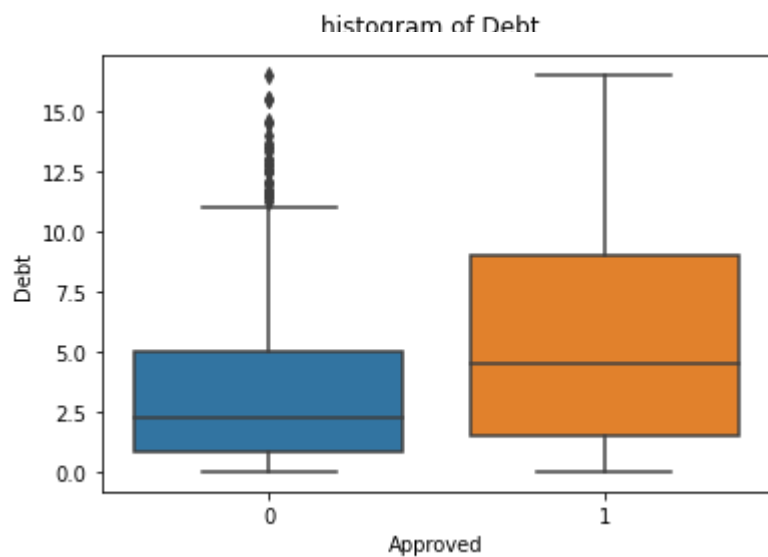




CreditScore:



Debt:



Statistical Testing using Anova:

In [57]:

```
from statsmodels.stats.multicomp import pairwise_tukeyhsd
from statsmodels.formula.api import ols
import statsmodels.api as smf
```

In [58]:

```
for i in continious:
    print("-----")
    print(i+":\n")
    print('ANOVA:\n')
    mod=ols(i+'~Approved',data=df).fit()
    aov_table=smf.stats.anova_lm(mod,type = 2)
    print(aov_table,'\n')
    print('Pvalue={}\n'.format(aov_table['PR(>F)'][0]))
    p=aov_table['PR(>F)'][0]

    if p>0.05:
        print(Fore.RED +"'{}' is a 'bad Predictor'\n".format(i))
        print('Avg of this feature is same for both card approved group and not approved gr
        print("p_val = {}\n".format(p))
    else:
        print('TUKEY:\n')
        print(Fore.RED +"'{}' is a 'good Predictor'\n".format(i))
        print('Avg of this feature is not same for both card approved group and not approve
        print('we need to perform Tuckey as atleast one category is different\n')
        print(Fore.GREEN +"'{}' is a 'good Predictor'\n".format(i))
        tukey=pairwise_tukeyhsd(df[i],df.Approved,alpha=0.05)
        print(tukey,'\n')
```

Income:

ANOVA:

	df	sum_sq	mean_sq	F	PR(>F)
Approved	1.0	4.161531e+06	4.161531e+06	104.814081	5.417547e-23
Residual	688.0	2.731631e+07	3.970393e+04	NaN	NaN

Pvalue=5.4175472814508e-23

TUKEY:

'Income' is a 'good Predictor'

Avg of this feature is not same for both card approved group and not approve
d group

we need to perform Tuckey as atleast one category is different

'Income' is a 'good Predictor'

Multiple Comparison of Means - Tukey HSD, FWER=0.05
=====

group1	group2	meandiff	p-adj	lower	upper	reject
0	1	156.2725	-0.0	126.3026	186.2424	True

Age:

ANOVA:

	df	sum_sq	mean_sq	F	PR(>F)
Approved	1.0	1565.625577	1565.625577	14.427262	0.000159
Residual	688.0	74660.762713	108.518550	NaN	NaN

Pvalue=0.00015856086530408587

TUKEY:

'Age' is a 'good Predictor'

Avg of this feature is not same for both card approved group and not approved group

we need to perform Tuckey as atleast one category is different

'Age' is a 'good Predictor'

Multiple Comparison of Means - Tukey HSD, FWER=0.05

```
=====
group1 group2 meandiff p-adj lower upper reject
-----
0      1      3.0311 0.0002 1.4643 4.5979 True
-----
```

YearsEmployed:

ANOVA:

	df	sum_sq	mean_sq	F	PR(>F)
Approved	1.0	200.03432	200.034320	100.012285	4.467613e-22
Residual	688.0	1376.06708	2.000098	NaN	NaN

Pvalue=4.467613399885831e-22

TUKEY:

'YearsEmployed' is a 'good Predictor'

Avg of this feature is not same for both card approved group and not approved group

we need to perform Tuckey as atleast one category is different

'YearsEmployed' is a 'good Predictor'

Multiple Comparison of Means - Tukey HSD, FWER=0.05

```
=====
group1 group2 meandiff p-adj lower upper reject
-----
0      1      1.0834 -0.0 0.8707 1.2962 True
-----
```

CreditScore:

ANOVA:

	df	sum_sq	mean_sq	F	PR(>F)
Approved	1.0	562.393969	562.393969	211.877834	4.856202e-42
Residual	688.0	1826.179944	2.654331	NaN	NaN

Pvalue=4.856202409405879e-42

TUKEY:

'CreditScore' is a 'good Predictor'

Avg of this feature is not same for both card approved group and not approved group

we need to perform Tuckey as atleast one category is different

'CreditScore' is a 'good Predictor'

Multiple Comparison of Means - Tukey HSD, FWER=0.05

```
=====
group1 group2 meandiff p-adj lower upper reject
-----
0      1      1.8167 -0.0 1.5716 2.0617  True
-----
```

Debt:

ANOVA:

	df	sum_sq	mean_sq	F	PR(>F)
Approved	1.0	469.578045	469.578045	28.0154	1.620321e-07
Residual	688.0	11531.860794	16.761426	NaN	NaN

Pvalue=1.6203209859251425e-07

TUKEY:

'Debt' is a 'good Predictor'

Avg of this feature is not same for both card approved group and not approved group

we need to perform Tuckey as atleast one category is different

'Debt' is a 'good Predictor'

Multiple Comparison of Means - Tukey HSD, FWER=0.05

```
=====
group1 group2 meandiff p-adj lower upper reject
-----
0      1      1.66    0.0 1.0442 2.2758  True
-----
```

In [59]:

```
df_industry = pd.get_dummies(df['Industry'], drop_first=True, prefix = "Industry")
```


In [61]:

```
df=pd.concat([df,df_industry],axis=1)
```

In [63]:

```
del df['Industry']
```

In [65]:

```
df_ethnicity = pd.get_dummies(df['Ethnicity'], drop_first=True, prefix = "Ethnicity")
```

In [67]:

```
df=pd.concat([df,df_ethnicity],axis=1)
```

In [68]:

```
del df['Ethnicity']
```

In [69]:

```
df_citizen = pd.get_dummies(df['Citizen'], drop_first=True, prefix = "Citizen")
```

In [71]:

```
df=pd.concat([df,df_citizen],axis=1)
```

In [74]:

```
del df['Citizen']
```

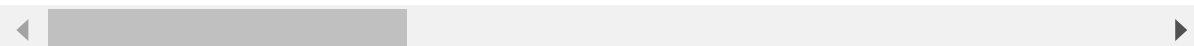
In [75]:

```
df.head()
```

Out[75]:

	Age	Debt	Married	BankCustomer	YearsEmployed	PriorDefault	Employed	CreditScore
0	30.83	0.000	1	1	1.25	1	1	1.0
1	58.67	4.460	1	1	3.04	1	1	6.0
2	24.50	0.500	1	1	1.50	1	0	0.0
3	27.83	1.540	1	1	3.75	1	1	5.0
4	20.17	5.625	1	1	1.71	1	0	0.0

5 rows × 30 columns



VIF Calculation

In [76]:

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

In [77]:

```
x = df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']]
```

In [78]:

```
x
```

Out[78]:

	Age	Debt	YearsEmployed	CreditScore	Income
0	30.83	0.000	1.25	1.0	0.0
1	58.67	4.460	3.04	6.0	560.0
2	24.50	0.500	1.50	0.0	824.0
3	27.83	1.540	3.75	5.0	3.0
4	20.17	5.625	1.71	0.0	0.0
...
685	21.08	10.085	1.25	0.0	0.0
686	22.67	0.750	2.00	2.0	394.0
687	25.25	13.500	2.00	1.0	1.0
688	17.92	0.205	0.04	0.0	750.0
689	35.00	3.375	2.46	0.0	0.0

690 rows × 5 columns

In [79]:

```
vif_data=pd.DataFrame()
```

In [81]:

```
vif_data['features']=x.columns
```

In [84]:

```
vif_data['VIF']=[variance_inflation_factor(x.values,i) for i in range (len(x.columns))]
```

In [85]:

```
vif_data
```

Out[85]:

	features	VIF
0	Age	2.828930
1	Debt	2.054744
2	YearsEmployed	2.298794
3	CreditScore	1.817621
4	Income	1.680393

Creating model

In [86]:

```
x=df.iloc[:,df.columns != 'Approved']
```

In [88]:

```
y=df.iloc[:,df.columns == 'Approved']
```

In [90]:

```
from sklearn.model_selection import train_test_split
```

In [91]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=123)
```

In [92]:

```
from sklearn.linear_model import LogisticRegression
```

In [93]:

```
model=LogisticRegression(solver='liblinear',random_state=123)
```

In [94]:

```
clf=model.fit(x_train,y_train)
```

C:\Conda dist\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

In [95]:

```
y_pred=clf.predict(x_test)
```

In [96]:

```
clf.score(x_test,y_test)
```

Out[96]:

```
0.8599033816425121
```

In [97]:

```
clf.intercept_
```

Out[97]:

```
array([-2.08619225])
```

In [98]:

```
clf.coef_
```

Out[98]:

```
array([[ -0.01793299, -0.02983195,  0.01303626,  0.58230414,  0.27553144,
         2.93605989,  0.28787561,  0.19055028, -0.00326723,  0.00294504,
        -1.30864209, -0.51328852,  0.38796879, -0.13284517, -0.64602139,
        -0.99670274,  0.05582364,  1.29475446, -0.05681709, -0.69727162,
        -0.25373214,  0.06279072,  0.89018486,  0.08004152, -0.22185161,
        -0.14085112, -0.37031655,  0.56010133,  1.01904202]])
```

Creating confusion matrix

In [99]:

```
from sklearn import metrics
```

In [103]:

```
cm=metrics.confusion_matrix(y_test,y_pred)
print(cm)
```

```
[[98 11]
 [18 80]]
```

In [104]:

```
cr=metrics.classification_report(y_test,y_pred)
print(cr)
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	109
1	0.88	0.82	0.85	98
accuracy			0.86	207
macro avg	0.86	0.86	0.86	207
weighted avg	0.86	0.86	0.86	207

ROC curve

In [107]:

```
from sklearn.metrics import roc_curve, roc_auc_score
```

In [110]:

```
y_pred_prob=clf.predict_proba(x_test)
print(y_pred_prob)
```

```
[0.98187248 0.01812752]
[0.21990199 0.78009801]
[0.83120417 0.16879583]
[0.59049407 0.40950593]
[0.27374389 0.72625611]
[0.13169043 0.86830957]
[0.23676165 0.76323835]
[0.05772132 0.94227868]
[0.06839222 0.93160778]
[0.91463342 0.08536658]
[0.05748357 0.94251643]
[0.14994097 0.85005903]
[0.2397884  0.7602116 ]
[0.82759593 0.17240407]
[0.44878724 0.55121276]
[0.06514231 0.93485769]
[0.23688485 0.76311515]
[0.49059855 0.50940145]
[0.15915356 0.84084644]
[0.62689163 0.37310837]
```

In [117]:

```
fpr,tpr,threshold = roc_curve(y_test,y_pred_prob[:,1])
```

In [112]:

```
roc_auc=roc_auc_score(y_test,y_pred_prob[:,1])
```

In [113]:

```
from sklearn.metrics import log_loss
```

In [114]:

```
log_loss(y_test,y_pred)
```

Out[114]:

4.838808265817111

In [119]:

```
plt.title('ROC curve for logloss:liblinear')
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.plot([0,1],[0,1],"r--")
plt.xlim([0,1])
plt.ylim([0,1])
plt.plot(fpr,tpr,label='AUC =' +str(roc_auc))
plt.legend(loc='best')
plt.show()
```

